

February 26, 2023

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[2]: df= pd.read_csv("Breast_cancer_data.csv")
```

```
[3]: df
```

```
[3]:      mean_radius  mean_texture  mean_perimeter  mean_area  mean_smoothness  \
0          17.99         10.38         122.80       1001.0         0.11840
1          20.57         17.77         132.90       1326.0         0.08474
2          19.69         21.25         130.00       1203.0         0.10960
3          11.42         20.38          77.58        386.1         0.14250
4          20.29         14.34         135.10       1297.0         0.10030
..          ...          ...          ...          ...          ...
564         21.56         22.39         142.00       1479.0         0.11100
565         20.13         28.25         131.20       1261.0         0.09780
566         16.60         28.08         108.30        858.1         0.08455
567         20.60         29.33         140.10       1265.0         0.11780
568          7.76         24.54          47.92        181.0         0.05263
```

```
      diagnosis
0            0
1            0
2            0
3            0
4            0
..          ...
564          0
565          0
566          0
567          0
568          1
```

```
[569 rows x 6 columns]
```

```
[5]: data= df.to_numpy()
X= data[:, 0:5]
print(X.shape)
y= data[:, 5]
m= y.shape[0]
y= y.reshape(m, 1)
X= np.insert(X, 0, np.ones((1,m)), axis= 1)
print(X)
```

```
(569, 5)
[[1.000e+00  1.799e+01  1.038e+01  1.228e+02  1.001e+03  1.184e-01]
 [1.000e+00  2.057e+01  1.777e+01  1.329e+02  1.326e+03  8.474e-02]
 [1.000e+00  1.969e+01  2.125e+01  1.300e+02  1.203e+03  1.096e-01]
 ...
 [1.000e+00  1.660e+01  2.808e+01  1.083e+02  8.581e+02  8.455e-02]
 [1.000e+00  2.060e+01  2.933e+01  1.401e+02  1.265e+03  1.178e-01]
 [1.000e+00  7.760e+00  2.454e+01  4.792e+01  1.810e+02  5.263e-02]]
```

```
[7]: from sklearn.model_selection import train_test_split
```

```
[11]: def hypothesis(X, theta):
    tmp = X@theta
    yest= 1/(1+(np.exp(-(tmp))))
    return yest

def cost(X, y, theta):
    m= X.shape[0]
    Yest= hypothesis(X, theta)
    lh1= np.log(Yest)
    lh2= np.log(1-Yest)
    cs= -((lh1.T@y)+lh2.T@(1-y))/m
    return cs

def gradient(X, y, theta):
    gr= np.zeros((6,1))
    Yest= hypothesis(X, theta)
    err= Yest- y
    m= X.shape[0]
    gr= X.T@err/m
    return gr

def normalize(X):
    cols= X.shape[1]
    Xmean= np.mean(X, axis= 0)
    Xmin= np.min(X, axis= 0)
    Xmax= np.max(X, axis= 0)
    X_norm = X.copy()
```

```

for i in range(1, cols):
    X_norm[:, i] = (X_norm[:, i] - Xmean[i]) / (Xmax[i] - Xmin[i])
return X_norm

```

```
X_norm = normalize(X)
```

```
[12]: X_norm
```

```

[12]: array([[ 1.          ,  0.18281548, -0.30130702,  0.21305346,  0.14681268,
              0.1989683 ],
             [ 1.          ,  0.30492254, -0.05139156,  0.28284822,  0.28467058,
             -0.10490459],
             [ 1.          ,  0.26327362,  0.06629528,  0.26280814,  0.23249667,
              0.11952441],
             ...,
             [ 1.          ,  0.11702912,  0.29727262,  0.11285306,  0.08619762,
             -0.10661985],
             [ 1.          ,  0.30634239,  0.3395452 ,  0.33260291,  0.25879571,
              0.19355167],
             [ 1.          , -0.3013532 ,  0.1775567 , -0.30439523, -0.20101341,
             -0.39478452]])

```

```

[13]: X_train, X_test, Y_train, Y_test = train_test_split(X_norm, y, test_size=0.2,
    ↪ random_state=0)
print("Number of training examples in train_set and test_set", X_train.
    ↪ shape[0], Y_train.shape[0])

```

Number of training examples in train_set and test_set 455 455

```

[ ]: N = 30000
alpha = 0.03
m = Y_train.shape[0]
theta = np.zeros((6, 1))
prev_cost = cost(X_train, Y_train, theta)
print("Cost before training:", prev_cost)
for i in range(N):
    print(i, prev_cost) # -----> J_history
    theta = theta - alpha * gradient(X_train, Y_train, theta)
    current_cost = cost(X_train, Y_train, theta)
    if abs(prev_cost - current_cost) < 1e-6:
        print(i)
        break
    prev_cost = current_cost
print("Final theta", theta)
print("\n")
print("Cost after Training:", prev_cost)

```

```
[22]: def findMSE(X, y, theta):
        m= X.shape[0]
        Y_pred= hypothesis(X, theta)
        Y_pred= [1 if i>= 0.5 else 0 for i in Y_pred]
        Y_pred= np.array(Y_pred).reshape(m, 1)
        err= Y_pred- y
        mse= err.T@err/m
        return mse

print("Train_mse:", findMSE(X_train, Y_train, theta))
```

Train_mse: [[0.07692308]]

```
[23]: print("Test_mse:", findMSE(X_test, Y_test, theta))
```

Test_mse: [[0.0877193]]

1 Confusion Matrix Evaluation

```
[24]: true_positive, true_negative, false_positive, false_negative= 0,0, 0, 0
m= X_test.shape[0]
Y_pred= hypothesis(X_test, theta)
Y_pred= [1 if i>= 0.5 else 0 for i in Y_pred]
Y_pred= np.array(Y_pred).reshape(m, 1)
for i in range(m):
    if Y_pred[i]== 0 and Y_test[i]==0:
        true_negative+=1
    elif Y_pred[i]==1 and Y_test[i]==0:
        false_positive+=1
    elif Y_pred[i]==1 and Y_test[i]==1:
        true_positive+=1
    elif Y_pred[i]==0 and Y_test[i]==1:
        false_negative+=1
print(true_positive, true_negative, false_positive, false_negative)
P = true_positive/(true_positive+false_positive)
R = true_positive/(true_positive+false_negative)
F1 = 2*P*R/(P+R)
print("Accuracy:",(true_positive+true_negative)/m)
print("P R F1",P,R,F1)
```

63 41 6 4

Accuracy: 0.9122807017543859

P R F1 0.9130434782608695 0.9402985074626866 0.9264705882352942

2 Logistic Regression on mnist dataset

```
[3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[4]: traindata= pd.read_csv("mnist_train.csv")
testdata= pd.read_csv("mnist_test.csv")
traindataset= traindata.to_numpy()
testdataset= testdata.to_numpy()
print(traindataset.shape, testdataset.shape)
train_labels= traindataset[:, 0].reshape(60000, 1)      # 60000 training samples
    ↳are present in traindataset
test_labels= testdataset[:, 0].reshape(10000, 1)        # 10000 testing
    ↳samples are present in testdataset
train_features= traindataset[:, 1:]
test_features= testdataset[:, 1:]
```

(60000, 785) (10000, 785)

3 sample Representation


```
[7]: i= 5      # ith sample in the data
np.printoptions(linewidth= 600)
print(train_features[i, :].reshape((28, 28)))           #the each samples have 784
    ↳values which can be reshaped into 28*28 matrix form
plt.imshow(train_features[i, :].reshape((28, 28)))      # imshow shows the pixels
    ↳by taking the values as the value of intensity
plt.show()
print("This is:", train_labels[i, 0])
```

```
[[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  13  25 100
 122  7  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  33 151 208 252 252
 252 146  0  0  0  0  0  0  0]]
```

```

[ 0 0 0 0 0 0 0 0 0 0 0 40 152 244 252 253 224 211
252 232 40 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 15 152 239 252 252 252 216 31 37
252 252 60 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 96 252 252 252 252 217 29 0 37
252 252 60 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 181 252 252 220 167 30 0 0 77
252 252 60 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 26 128 58 22 0 0 0 0 100
252 252 60 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 157
252 252 60 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 110 121 122 121 202
252 194 3 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 10 53 179 253 253 255 253 253
228 35 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 5 54 227 252 243 228 170 242 252 252
231 117 6 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 6 78 252 252 125 59 0 18 208 252 252
252 252 87 7 0 0 0 0 0 0]
[ 0 0 0 0 0 0 5 135 252 252 180 16 0 21 203 253 247 129
173 252 252 184 66 49 49 0 0 0]
[ 0 0 0 0 0 3 136 252 241 106 17 0 53 200 252 216 65 0
14 72 163 241 252 252 223 0 0 0]
[ 0 0 0 0 0 105 252 242 88 18 73 170 244 252 126 29 0 0
0 0 0 89 180 180 37 0 0 0]
[ 0 0 0 0 0 231 252 245 205 216 252 252 252 124 3 0 0 0
0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 207 252 252 252 252 178 116 36 4 0 0 0 0
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 13 93 143 121 23 6 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0]]

```



This is: 2

4 Normalize features

```
[8]: # Here normalization is done by dividing the dataset with max value in the data,
      ↪ i.e in MNIST the max value is approx 255.
      train_features= train_features/255
      test_features= test_features/255
```

```
[9]: m_train = train_features.shape[0]
      m_test= test_features.shape[0]
      train_features= np.insert(train_features, 0, np.ones((1, m_train)), axis= 1)
      test_features= np.insert(test_features, 0, np.ones((1, m_test)), axis= 1)
```

```
[14]: def hypothesis(X, theta):
        tmp = X@theta
        Yest= 1/(1+(np.exp(-tmp)))
        return Yest

      def cost(X, y, theta):
        m= X.shape[0]
```

```

    lh1= np.log(hypothesis(X, theta))
    lh2= np.log(1- (hypothesis(X, theta)))
    cs= -(lh1.T@y+lh2.T@(1-y))/m
    return cs
def gradient(X, y, theta):
    m= y.shape[0]
    Yest= hypothesis(X, theta)
    err= Yest- y
    gr= (X.T@err)/m
    return gr

train_labels0= np.array([1 if i==0 else 0 for i in train_labels]).
    ↪reshape(m_train, 1)
train_labels1= np.array([1 if i==1 else 0 for i in train_labels]).
    ↪reshape(m_train, 1)
train_labels2= np.array([1 if i==2 else 0 for i in train_labels]).
    ↪reshape(m_train, 1)
train_labels3= np.array([1 if i==3 else 0 for i in train_labels]).
    ↪reshape(m_train, 1)
train_labels4= np.array([1 if i==4 else 0 for i in train_labels]).
    ↪reshape(m_train, 1)
train_labels5= np.array([1 if i==5 else 0 for i in train_labels]).
    ↪reshape(m_train, 1)
train_labels6= np.array([1 if i==6 else 0 for i in train_labels]).
    ↪reshape(m_train, 1)
train_labels7= np.array([1 if i==7 else 0 for i in train_labels]).
    ↪reshape(m_train, 1)
train_labels8= np.array([1 if i==8 else 0 for i in train_labels]).
    ↪reshape(m_train, 1)
train_labels9= np.array([1 if i==9 else 0 for i in train_labels]).
    ↪reshape(m_train, 1)

```

```

[15]: def gradient_descent(X, y, theta_initial, alpha, tol, N):
    theta= theta_initial
    prev_cost= cost(X, y, theta)
    print("Cost before training:", prev_cost)
    for i in range(N):
        print("-", end= "")
        theta= theta- alpha*gradient(X, y, theta)
        cs= cost(X, y, theta)
        if abs(prev_cost- cs)<tol:
            print(i)
            break
        prev_cost= cs
    print("the cost after training:", prev_cost)
    return theta

```



```
[16]: theta_initial= np.zeros((785, 1))
      theta0= gradient_descent(train_features, train_labels0, theta_initial, 0.03,
      ↪1e-6, 3000 )
      theta1= gradient_descent(train_features, train_labels1, theta_initial, 0.03,
      ↪1e-6, 3000 )
      theta2= gradient_descent(train_features, train_labels2, theta_initial, 0.03,
      ↪1e-6, 3000 )
      theta3= gradient_descent(train_features, train_labels3, theta_initial, 0.03,
      ↪1e-6, 3000 )
      theta4= gradient_descent(train_features, train_labels4, theta_initial, 0.03,
      ↪1e-6, 3000 )
      theta5= gradient_descent(train_features, train_labels5, theta_initial, 0.03,
      ↪1e-6, 3000 )
      theta6= gradient_descent(train_features, train_labels6, theta_initial, 0.03,
      ↪1e-6, 3000 )
      theta7= gradient_descent(train_features, train_labels7, theta_initial, 0.03,
      ↪1e-6, 3000 )
      theta8= gradient_descent(train_features, train_labels8, theta_initial, 0.03,
      ↪1e-6, 3000 )
      theta9= gradient_descent(train_features, train_labels9, theta_initial, 0.03,
      ↪1e-6, 3000 )
```

Cost before training: `[[0.69314718]]`

[illegible]

```
-----the cost after training: [[0.04231707]]
Cost before training: [[0.69314718]]
-----
```

-----the cost after training: [[0.04047849]]
Cost before training: [[0.69314718]]

-----the cost after training: [[0.08700877]]
Cost before training: [[0.69314718]]

-----the cost after training: [[0.10102085]]
Cost before training: [[0.69314718]]

-----the cost after training: [[0.07153784]]
Cost before training: [[0.69314718]]

```
-----the cost after training: [[0.1134967]]
Cost before training: [[0.69314718]]
-----
```

```
-----the cost after training: [[0.05408988]]
Cost before training: [[0.69314718]]
```

```
-----the cost after training: [[0.06173994]]
Cost before training: [[0.69314718]]
```

-----the cost after training: [[0.15234157]]
Cost before training: [[0.69314718]]


```
-----the cost after training: [[0.12399921]]
```

```
[18]: theta= np.concatenate((theta0, theta1, theta2, theta3, theta4, theta5, theta6,
    ↪ theta7, theta8, theta9), axis= 1)
```

```
[20]: theta.shape
```

[20]: (785, 10)

```
[21]: theta
```

```
[21]: array([[ -0.67947939, -0.12596979, -0.76482933, ..., -0.3033674 ,
          -1.68213952, -0.9087435 ],
          [ 0.          , 0.          , 0.          , ..., 0.          ,
            0.          , 0.          ],
          [ 0.          , 0.          , 0.          , ..., 0.          ,
            0.          , 0.          ],
          ...,
          [ 0.          , 0.          , 0.          , ..., 0.          ,
```

```

0.      , 0.      ],
[ 0.      , 0.      , 0.      , ..., 0.      ,
 0.      , 0.      ],
[ 0.      , 0.      , 0.      , ..., 0.      ,
 0.      , 0.      ]])

```

```

[25]: testPrediction = 1/(1+np.exp(-test_features@theta))
      i = 1003
      print(testPrediction[i,:])
      print(np.argmax(testPrediction[i,:]))
      print(test_labels[i])

```

```

[3.61221678e-04 3.14821596e-07 5.16941061e-04 4.00406839e-01
 1.89973454e-06 5.53674883e-01 3.82089047e-07 2.16860535e-06
 9.53381239e-02 2.06206713e-05]
5
[5]

```

```

[26]: testPred = np.argmax(testPrediction,axis=1).reshape(10000,1)
      print(testPred)

```

```

[[7]
 [2]
 [1]
 ...
 [4]
 [8]
 [6]]

```

```

[27]: correct = [1 if testPred[i]==test_labels[i] else 0 for i in range(0,10000)]
      accuracy = np.sum(correct)/10000
      print(accuracy*100)

```

89.92

5 Logistic_regression on university admission dataset

```

[1]: import pandas as pd
      import numpy as np
      import seaborn as sns
      import matplotlib.pyplot as plt
      %matplotlib inline

```

```

[2]: df = pd.read_csv('E:\VS_
      ↪CODE\machine-learning-andrew-ng-master\machine-learning-andrew-ng-master\data\ex2data1.
      ↪txt', sep=',', header=None)
      df.columns = ['exam_score_1', 'exam_score_2', 'label']

```

```
[3]: df
```

```
[3]:    exam_score_1  exam_score_2  label
0      34.623660    78.024693      0
1      30.286711    43.894998      0
2      35.847409    72.902198      0
3      60.182599    86.308552      1
4      79.032736    75.344376      1
..      ...      ...      ...
95     83.489163    48.380286      1
96     42.261701    87.103851      1
97     99.315009    68.775409      1
98     55.340018    64.931938      1
99     74.775893    89.529813      1
```

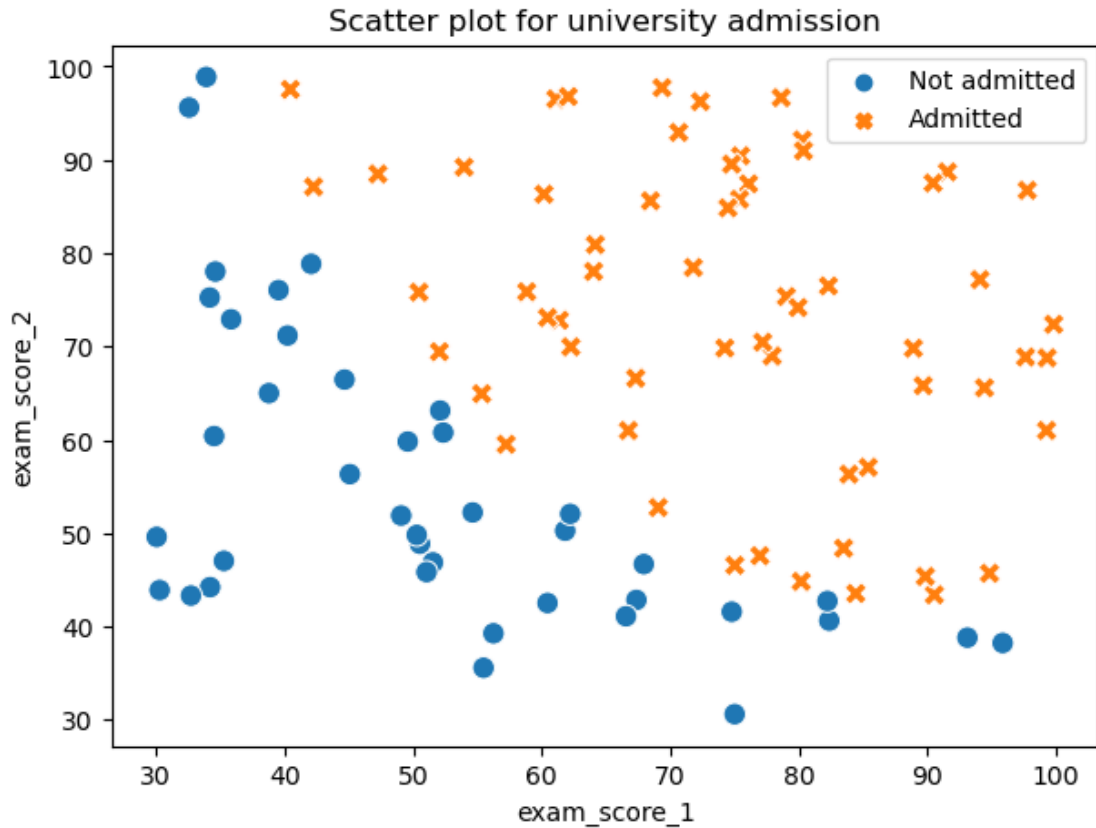
[100 rows x 3 columns]

```
[4]: df.describe().T
```

```
[4]:    count      mean      std      min      25%      50%  \
exam_score_1  100.0  65.644274  19.458222  30.058822  50.919511  67.032988
exam_score_2  100.0  66.221998  18.582783  30.603263  48.179205  67.682381
label         100.0   0.600000   0.492366   0.000000   0.000000   1.000000

      75%      max
exam_score_1  80.212529  99.827858
exam_score_2  79.360605  98.869436
label         1.000000   1.000000
```

```
[8]: plt.figure(figsize=(7,5))
plot= sns.scatterplot(x= "exam_score_1", y= "exam_score_2", data= df, hue=
    ↪ "label", style= "label", s= 80)
#hue---> here hue means how color to be differentiated, so here labe column is
    ↪ used to differentiate the color
#style---> Here style means how to differentiate the markers to be used, so
    ↪ here the lanel column is used to differentiate the style
# s--> size of marker
handles, labels= plot.get_legend_handles_labels()
plot.legend(handles[0:], ["Not admitted", "Admitted"])
plt.title("Scatter plot for university admission")
plt.show()
```



```
[92]: def sigmoid(z):
        z= np.array(z)
        return 1/(1+np.exp(-z))
```

```
[123]: def cost_function(theta, X, y):
        m = y.shape[0]
        theta = theta[:, np.newaxis] #trick to make numpy minimize work
        h = sigmoid(X.dot(theta))
        J = (1/m) * (-y.T.dot(np.log(h)) - (1-y).T.dot(np.log(1-h)))

        diff_hy = h - y
        grad = (1/m) * diff_hy.T.dot(X)

        return J, grad
```

```
[124]: m = df.shape[0]
        X = np.hstack((np.ones((m,1)),df[['exam_score_1', 'exam_score_2']].values))
        y = np.array(df.label.values).reshape(-1,1)
        initial_theta = np.zeros(shape=(X.shape[1]))
```

```
[125]: cost, grad = cost_function(initial_theta, X, y)
print('Cost at initial theta (zeros):', cost)
print('Expected cost (approx): 0.693')
print('Gradient at initial theta (zeros):')
print(grad.T)
print('Expected gradients (approx):\n -0.1000\n -12.0092\n -11.2628')
```

```
Cost at initial theta (zeros): [[0.69314718]]
Expected cost (approx): 0.693
Gradient at initial theta (zeros):
[[ -0.1      ]
 [-12.00921659]
 [-11.26284221]]
Expected gradients (approx):
-0.1000
-12.0092
-11.2628
```

```
[126]: import scipy.optimize as opt
def optimize_theta(X, y, initial_theta):
    opt_results = opt.minimize(cost_function, initial_theta, args=(X, y),
    ↪method='TNC',
                                jac=True, options={'maxiter':400})
    return opt_results['x'], opt_results['fun']
```

```
[127]: opt_theta, cost = optimize_theta(X, y, initial_theta)
```

```
C:\Users\JAGADISH\AppData\Local\Temp\ipykernel_10980\1274892550.py:3:
DeprecationWarning: 'maxiter' has been deprecated in favor of 'maxfun' and will
be removed in SciPy 1.11.0.
    opt_results = opt.minimize(cost_function, initial_theta, args=(X, y),
method='TNC',
```

```
[128]: opt_theta
```

```
[128]: array([-25.16131862,  0.20623159,  0.20147149])
```

```
[129]: cost
```

```
[129]: 0.20349770158947464
```

6 Decision Boundary

```
[156]: plt.figure(figsize= (7, 5))
plot= sns.scatterplot(x= "exam_score_1", y= "exam_score_2", data= df, hue=
    ↪"label", style= "label", s= 80)
```

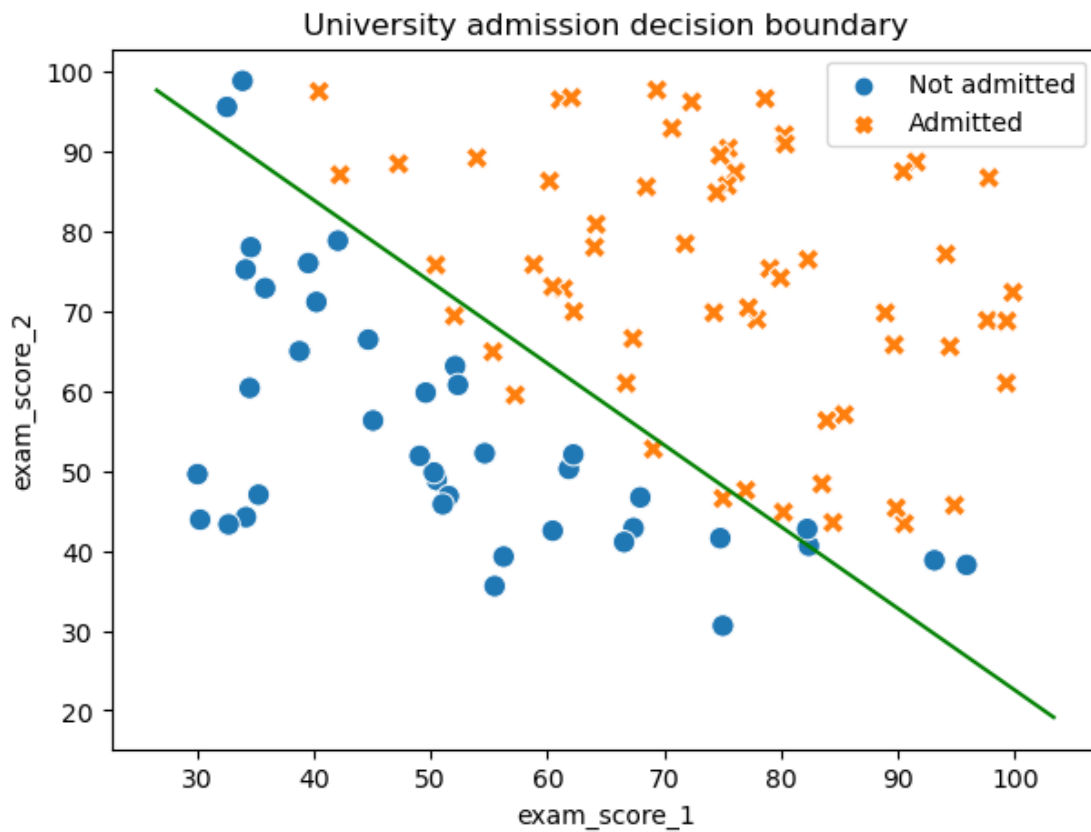
```

handles, labels= plot.get_legend_handles_labels()
plot.legend(handles[0:], ["Not admitted", "Admitted"])
plt.title("University admission decision boundary")

plot_x= np.array(plot.get_xlim())
print(plot_x)
plot_y = (-opt_theta[1]/opt_theta[2]) * plot_x - opt_theta[0]/opt_theta[2]
plt.plot(plot_x, plot_y, "--", c= "green")
plt.show(plot)

```

```
[ 26.57037068 103.31630956]
```



7 Evaluating Logistic Regression

```

[157]: prob= sigmoid(np.array([1, 45, 85]).dot(opt_theta))
print("The prob of getting admission of student with 45 and 85 marks is:", prob)

```

```

The prob of getting admission of student with 45 and 85 marks is:
0.7762906236225744

```

8 Accuracy on training set

```
[160]: def predict(X, theta):  
        y_pred= [1 if sigmoid(X[i, :].dot(theta))>= 0.5 else 0 for i in range(0, X.  
        ↪shape[0])]   
        return y_pred
```

```
[161]: y_pred_prob = predict(X, opt_theta)
```

```
[162]: y_pred_prob
```

```
[162]: [0,  
        0,  
        0,  
        1,  
        1,  
        0,  
        1,  
        0,  
        1,  
        1,  
        1,  
        0,  
        1,  
        1,  
        0,  
        1,  
        0,  
        1,  
        0,  
        0,  
        1,  
        1,  
        0,  
        1,  
        0,  
        0,  
        1,  
        1,  
        1,  
        1,  
        1,  
        0,  
        0,  
        1,  
        1,  
        0,  
        0,  
        0,  
        0,  
        0,
```

1,
1,
0,
0,
1,
0,
1,
1,
0,
0,
1,
1,
1,
1,
1,
1,
1,
1,
0,
0,
0,
0,
1,
1,
1,
1,
1,
1,
0,
0,
0,
0,
0,
1,
0,
1,
1,
1,
0,
1,
1,
1,
1,
1,
1,
1,
1,
1,
0,
1,
1,
1,
1,


```
0,  
1,  
1,  
0,  
1,  
1,  
0,  
1,  
1,  
0,  
1,  
1,  
1,  
1,  
1,  
0,  
1]
```

```
[165]: y_pred_prob== df.label.values
```

```
[165]: array([ True,  True,  True,  True,  True,  True,  True, False,  True,  
         True, False,  True,  True,  True,  True,  True, False,  True,  
         True,  True,  True,  True,  True,  True,  True,  True,  True,  
        False,  True,  True,  True,  True,  True,  True, False,  True,  True,  
        False,  True,  True,  True,  True,  True,  True, False,  True,  
         True,  True,  True,  True,  True,  True,  True,  True,  True,  
         True,  True,  True, False,  True,  True,  True,  True,  True,  
         True,  True,  True,  True,  True,  True,  True,  True,  True,  
         True,  True,  True,  True,  True,  True,  True, False,  True,  
         True,  True, False,  True,  True,  True,  True,  True,  True,  
         True,  True,  True,  True,  True,  True,  True,  True, False,  
         True])
```

```
[163]: f'Train Accuracy: {np.mean(y_pred_prob== df.label.values)*100}'
```

```
[163]: 'Train Accuracy: 89.0'
```

9 Equivalent with Sklearn

```
[167]: from sklearn.linear_model import LogisticRegression  
log_reg= LogisticRegression(solver= "newton-cg", max_iter= 400)  
log_reg.fit(df[['exam_score_1', "exam_score_2"]].values, df.label.values)
```

```
[167]: LogisticRegression(max_iter=400, solver='newton-cg')
```

```
[168]: log_reg.intercept_, log_reg.coef_
```

```
[168]: (array([-25.05200378]), array([[0.2053533 , 0.20058239]]))
```

#Accuracy with sklearn

```
[170]: log_reg.score(df[['exam_score_1', 'exam_score_2']].values,df.label.values)
```

```
[170]: 0.89
```

10 Regularized Logistic Regression

Visualization of the data

```
[181]: df2= pd.read_csv("E:\VS_
↳CODE\machine-learning-andrew-ng-master\machine-learning-andrew-ng-master\data\ex2data2.
↳txt", sep =",", header= None)
df2.columns= ['test_1', 'test_2', 'label']
```

```
[182]: df2
```

```
[182]:
```

	test_1	test_2	label
0	0.051267	0.699560	1
1	-0.092742	0.684940	1
2	-0.213710	0.692250	1
3	-0.375000	0.502190	1
4	-0.513250	0.465640	1
..
113	-0.720620	0.538740	0
114	-0.593890	0.494880	0
115	-0.484450	0.999270	0
116	-0.006336	0.999270	0
117	0.632650	-0.030612	0

[118 rows x 3 columns]


```
[183]: df2.describe().T
```

```
[183]:
```

	count	mean	std	min	25%	50%	75%	\
test_1	118.0	0.054779	0.496654	-0.83007	-0.372120	-0.006336	0.478970	
test_2	118.0	0.183102	0.519743	-0.76974	-0.254385	0.213455	0.646563	
label	118.0	0.491525	0.502060	0.00000	0.000000	0.000000	1.000000	

	max
test_1	1.0709
test_2	1.1089
label	1.0000

```
[184]: plt.figure(figsize=(7,5))
plot= sns.scatterplot(x= "test_1", y= "test_2", data= df2, hue= "label", style=
↪"label", s= 80)
handles, labels= plot.get_legend_handles_labels()
plt.legend(handles[0:], ["Accepted", "Rejected"])
plt.title("Scatter plot for the exam data")
plt.show()
```



```
[187]: def map_feature(X1, X2, degree):
X1 = np.array(X1).reshape(-1,1)
X2 = np.array(X2).reshape(-1,1)

out = np.ones((X1.shape[0], 1))
for i in range(1, degree+1):
    for j in range(0, i+1):
        p = (X1**(i-j)) * (X2**j)
        out = np.append(out, p, axis=1)
return out
```

```
[188]: X_p = map_feature(df2.test_1.values, df2.test_2.values, 6)
X_p.shape
```

```
[188]: (118, 28)
```

```
[189]: def cost_function_reg(theta, X, y, lambda_reg):  
    m = y.shape[0]  
    theta = theta[:, np.newaxis]  
    h = sigmoid(X.dot(theta))  
    J = (1/m) * (-y.T.dot(np.log(h)) - (1-y).T.dot(np.log(1-h))) + (lambda_reg/  
    ↪(2*m)) * np.sum(theta[1:])**2  
  
    diff_hy = h - y  
    grad = (1/m) * diff_hy.T.dot(X) + ((lambda_reg/m) * theta.T)  
    grad[0, 0] = (1/m) * diff_hy.T.dot(X[:, 0])  
  
    return J, grad
```

```
[190]: import scipy.optimize as opt  
def optimize_theta_reg(X, y, initial_theta, lambda_reg):  
    opt_results = opt.minimize(cost_function_reg, initial_theta, args=(X, y, ↵  
    ↪lambda_reg), method='TNC', jac=True, options={'maxiter':400})  
    return opt_results['x'], opt_results['fun']
```

```
[191]: m = df.shape[0]  
X = X_p  
y = np.array(df2.label.values).reshape(-1,1)  
initial_theta = np.zeros(shape=(X.shape[1]))
```

```
[192]: lambda_reg = 1  
cost, grad = cost_function_reg(initial_theta, X, y, lambda_reg)
```

```
[193]: print(grad.T[:5])
```

```
[[8.47457627e-03]  
 [1.87880932e-02]  
 [7.77711864e-05]  
 [5.03446395e-02]  
 [1.15013308e-02]]
```

```
[194]: lambda_reg = 10  
initial_theta = np.ones(shape=(X.shape[1]))  
cost, grad = cost_function_reg(initial_theta, X, y, lambda_reg)
```

```
[195]: print(grad.T[:5])
```

```
[[0.34604507]  
 [0.16135192]  
 [0.19479576]  
 [0.22686278]]
```

```
[0.09218568]]
```

```
[196]: lambda_reg = [1, 10, 100, 0]
fig, axs = plt.subplots(nrows=1, ncols=4, figsize=(15,4))
u = np.linspace(-1, 1.5, 50)
v = np.linspace(-1, 1.5, 50)

for il, l in enumerate(lambda_reg):
    theta_opt, cost = optimize_theta_reg(X, y, initial_theta, l)
    z = np.zeros((u.shape[0], v.shape[0]))
    for i in range(len(u)):
        for j in range(len(v)):
            z[i,j] = map_feature(u[i], v[j], 6).dot(theta_opt)

    sns.scatterplot(x='test_1', y='test_2', hue='label', data=df2,
                    style='label', s=80, ax=axs[il])

    axs[il].contour(u, v, z.T, levels=[0], colors='green')
    axs[il].set_title('$\lambda={}$'.format(l))
fig.tight_layout()
plt.show()
```

C:\Users\JAGADISH\AppData\Local\Temp\ipykernel_10980\968707659.py:3:

DeprecationWarning: 'maxiter' has been deprecated in favor of 'maxfun' and will be removed in SciPy 1.11.0.

```
opt_results = opt.minimize(cost_function_reg, initial_theta, args=(X, y,
lambda_reg), method='TNC', jac=True, options={'maxiter':400})
```

C:\Users\JAGADISH\AppData\Local\Temp\ipykernel_10980\968707659.py:3:

DeprecationWarning: 'maxiter' has been deprecated in favor of 'maxfun' and will be removed in SciPy 1.11.0.

```
opt_results = opt.minimize(cost_function_reg, initial_theta, args=(X, y,
lambda_reg), method='TNC', jac=True, options={'maxiter':400})
```

C:\Users\JAGADISH\AppData\Local\Temp\ipykernel_10980\968707659.py:3:

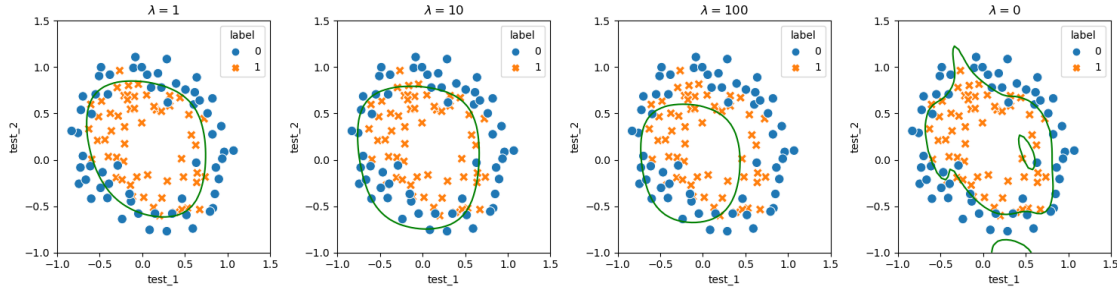
DeprecationWarning: 'maxiter' has been deprecated in favor of 'maxfun' and will be removed in SciPy 1.11.0.

```
opt_results = opt.minimize(cost_function_reg, initial_theta, args=(X, y,
lambda_reg), method='TNC', jac=True, options={'maxiter':400})
```

C:\Users\JAGADISH\AppData\Local\Temp\ipykernel_10980\968707659.py:3:

DeprecationWarning: 'maxiter' has been deprecated in favor of 'maxfun' and will be removed in SciPy 1.11.0.

```
opt_results = opt.minimize(cost_function_reg, initial_theta, args=(X, y,
lambda_reg), method='TNC', jac=True, options={'maxiter':400})
```



```
[197]: lambda_reg = 1
theta, cost = optimize_theta_reg(X, y, initial_theta, lambda_reg)
theta
```

C:\Users\JAGADISH\AppData\Local\Temp\ipykernel_10980\968707659.py:3:
 DeprecationWarning: 'maxiter' has been deprecated in favor of 'maxfun' and will
 be removed in SciPy 1.11.0.

```
opt_results = opt.minimize(cost_function_reg, initial_theta, args=(X, y,
lambda_reg), method='TNC', jac=True, options={'maxiter':400})
```

```
[197]: array([ 1.27273509,  0.62525435,  1.18108521, -2.01994882, -0.91742556,
-1.43167368,  0.12399628, -0.36552234, -0.35723208, -0.17514253,
-1.4581339 , -0.05098852, -0.61553085, -0.27470069, -1.19280263,
-0.24220871, -0.20601057, -0.04472767, -0.2777735 , -0.29536755,
-0.45637086, -1.04318579,  0.02776829, -0.29241701,  0.01556523,
-0.32737793, -0.14388044, -0.92463148])
```

```
[198]: y_pred_prob = predict(X, theta)
f'Train accuracy: {np.mean(y_pred_prob == df2.label.values) * 100}'
```

```
[198]: 'Train accuracy: 83.05084745762711'
```

```
[199]: from sklearn.linear_model import LogisticRegression
log_reg = LogisticRegression(solver='newton-cg', max_iter=400)
log_reg.fit(X[:,1:], df2.label.values)
```

```
[199]: LogisticRegression(max_iter=400, solver='newton-cg')
```

```
[200]: log_reg.intercept_, log_reg.coef_
```

```
[200]: (array([1.27273852]),
array([[ 0.62527427,  1.18107953, -2.01995701, -0.91743361, -1.43166228,
0.12400943, -0.36552879, -0.35723375, -0.1751281 , -1.45816817,
-0.05099315, -0.61556795, -0.27470949, -1.19281161, -0.24218951,
-0.20599958, -0.04473522, -0.27778736, -0.29537501, -0.45635027,
-1.04321271,  0.02777197, -0.29243756,  0.0155633 , -0.32738395,
```

```
-0.14388956, -0.92464266]]))
```

```
[201]: log_reg.score(X[:,1:], df2.label.values)
```

```
[201]: 0.8305084745762712
```

```
[ ]:
```